

Towards a Strongly Typed Functional Operating System*

Arjen van Weelden and Rinus Plasmeijer

Computer Science Institute
University of Nijmegen
Tooveld 1, 6525 ED Nijmegen, The Netherlands
`{arjenw,rinus}@cs.kun.nl`

In this paper, we present Famke. It is a prototype implementation of a strongly typed operating system written in Clean. Famke provides the creation and management of independent distributed Clean processes on a network of workstations. It uses Clean's dynamic type system and a dynamic linker to communicate values of any type, e.g. integers, floats and functions (i.e. compiled code), between running applications in a type safe way. Mobile processes can be implemented using Clean's ability to communicate functions. We have built an interactive shell for Famke that enables user interaction. The shell uses a command language that allows construction of new processes. Famke's type checks the command line before executing it. Famke's design and implementation makes it a strongly typed operating system tailored to a given situation.

n

Functional programming languages like Haskell [1] and Clean [2,3] offer a very powerful static type system. Compact, reusable, and readable programs can be written in these languages while the static type system is able to detect and report programming errors at compile time. But this works only within a

limited context. Developed applications often need to communicate with each other. Therefore the communication of objects to take place in a type safe way is not only for simple objects, but objects of any type, including functions. Therefore, this is not easy to realize: the compile time type information must be kept inside a compiled executable, and therefore cannot be

elden and Rinus Plasmeijer

is mainly because the used programming language has no
ngs.

re present a prototype implementation of a micro kernel,
(functional *micro kernel experiment*). It provides explicit
ncurrency and type safe message passing for all types to
Clean. By adding servers that provide common operating
entire strongly typed, distributed operating system can be
ke.

a powerful dynamic type system [4] for this purpose and a
extend a running application with new code. Fortunately,
n offers some of the required basic facilities: it offers a hybrid
tic as well as dynamic typing (*dynamics*) [5], including run-
dynamic linking [6] (currently on Microsoft Windows only).
communication, Famke uses the above mentioned facilities
implement *lightweight threads, processes, exception handling*
ge passing without requiring additional language constructs

se of an underlying operating system to avoid some low-
n work and to integrate better with existing software (e.g.
ne console and the file system). With Famke, we want to
ving objectives without changing the Clean compiler or run-

face (API) for Clean programmers with which it is easy to
(ed) processes that can communicate expressions of any type
ay;

active shell with which it is easy to manage, apply and com-
) processes, and even construct new processes interactively.
l type check the command line before executing it in order
arly;

lar design using an extensible micro kernel approach;
le system by using static types where possible and, if static
be done (e.g. between different programs), dynamic type

n that is easy to port to another operating system (if the
pports it).

in Clean

been extended with a polymorphic dynamic type system and its static type system. Here, we will give a small introduction to it. A dynamic is a value of type `Dynamic` which contains a representation of the type of that value.

is formed (i.e. lifted from the static to the dynamic type system) by the `dynamic` in combination with the value and an optional type pattern. The compiler will infer the type, separated by a double colon.

```
g | Just a2
```

```
c -> Maybe Int
  = Just x
  = Nothing
```

`Dynamic` can be matched in function alternatives and case expressions. It can be lifted from the dynamic back into the static type system. Such expressions consist of an optional value pattern and a type pattern. In the `matchInt` function, `Just` returns the value contained inside the dynamic and `Nothing` if it has any other type. The compiler translates these into run-time type unifications. If the unification fails, the next alternative is tried, as in a common pattern match.

```
dynamic Dynamic -> Dynamic3
a -> b) (x :: a) = dynamic f x :: b
          = dynamic "Error: cannot apply"
```

can contain type variables which, if the run-time unification succeeds, are unified with the offered type. In the example above, `dynamicApply` is a function `f` inside its first argument can be unified with the value `x` inside the second argument. If this is the case then we can safely apply `f` to `x`. The result of this application has type `b`. It is generally unknown what this type `b` will be. The result is a dynamic (and only a dynamic) again, because the type is not instantiated by the run-time unification.

```
dynamic -> Maybe t | TC t4
```

elden and Rinus Plasmeijer

variable in a pattern associates it with the type variable with the static type of the function. The static type variable then appears in the predefined TC (type code) class [5]. In the example type `t` will be determined by the static context in which it appears. It may impose a restriction on the actual type that is accepted at run-time. The type `Dynamic` is a special type that is accepted at run-time. It yields `Just` the value inside the dynamic (if the value of the required context dependent type) or `Nothing`.

The run-time system of Clean [6] supports writing dynamics to extend a program in again, possibly in another program or during another execution of the program.

```
type Dynamic *World -> (Bool, *World)5  
type *World -> (Bool, Dynamic, *World)
```

will be read in lazily after a successful run-time unification (pattern match on the dynamic). The amount of data and code that the linker will link in, is therefore determined by the amount of data and code inside the dynamic. Dynamics written by a program can be used by any other program, providing a form of persistence and a mechanism for communication.

Clean, as well as other functional languages, to construct new functions (including and higher-order functions) in combination with Clean's run-time linking, enables us to extend a running application. The extension can be type checked after which it is guaranteed to fit.

Famke

A programmer can construct concurrent programs in Clean, using process management and exception handling primitives.

Clean offers only very limited library support for process management and exception handling.

The concurrent Clean [7] did offer sophisticated support for parallel execution of lightweight processes, but no support for exception handling. It was targeted at deterministic, implicit concurrency, but we can also use it for non-deterministic, explicit concurrency. The concurrent Clean [7] did offer sophisticated support for parallel execution of lightweight processes, but no support for exception handling. It was targeted at deterministic, implicit concurrency, but we can also use it for non-deterministic, explicit concurrency.

very heavyweight, and it would prevent them from sharing the state. This still would not have exception handling. The kernel does her own scheduling of threads in order to keep them busy and provide exception handling.

Implementation

To support cooperative threads we need a way to suspend running threads and to resume them later. Wand [8] shows that this can be done using the call/CC construct offered by Scheme and other functional languages. We copy this approach using first class continuations. Because Clean has no call/CC construction, we have to write continuations explicitly. Our approach closely resembles Claessen's [9], but our primitives operate directly on the kernel state instead of using monads. We have extended the implementation with support for exception handling (see section 3.2).

```

cont (a -> KernelOp) -> KernelOp6
kernel -> Kernel

let cont a
  in let kernel -> cont x kernel '
    in let //7 calculate argument for cont
      in let kernel... // operate on the kernel state

```

The type `Thread`, such as the example function above, gets the kernel state (named `cont`; of type `a -> KernelOp`) as its argument and returns a new computation step, which calculates the argument of the next kernel computation, to form a new function (of type `KernelOp`). This function, when evaluated on a kernel state (named `kernel`; of type `Kernel`), returns a new kernel state.

```

abstract thread id

newThread :: ThreadId, newId :: ThreadId,
  world :: [ThreadState], world :: *World}

newThread :: ThreadId, thrCont :: KernelOp}

```

elden and Rinus Plasmeijer

ready list), and the `world` state which is provided by the Clean
language's uniqueness type system makes these types a little more
convenient. We will not show this in the examples in order to keep them

```
and a) -> Thread ThreadId
\cont k={newId, ready}9 ->
newId = inc newId, ready = [threadState:ready]}10

ThreadId = newId, thrCont = thread (\_ k -> k)}
```

```
ThreadId
: {currentId} -> cont currentId k
```

This function starts the given thread concurrently with the other
threads. It is evaluated for their effect on the kernel and the world state.
It does not return a result, hence the polymorphically parameterized
type. This relieves our system from the additional complexity of return-
ing a result to the parent thread. The communication primitives that will be
described enable programmers to extend the `newThread` primitive to de-
pend on the parent. Threads can obtain their thread identification with

the `threadId` function. The scheduling of threads is done cooperatively. This means that threads must
yield when they are scheduled, and should not run endless tight
loops. The `threadId` function then evaluates the next ready thread. `StartFamke`
is the standard Clean `Start` function to start the evaluation of

```
Kernel -> Kernel
currentId, ready} = {k & ready = ready ++ [threadState]}

ThreadId = currentId, thrCont = cont}
```

```
-> Kernel
= []} = k // nothing to schedule
= [{thrId, thrCont}:tail]} =
ready = tail, currentId = thrId}
k' // evaluate the thread until it yields
```

```

= [threadState], world = world}
thrId = firstId, thrCont = mainThread (\_ k -> k)}

```

is currently being evaluated returns directly to the scheduler as a `yield` action, because `yield` does not evaluate the tail. Instead, it stores the continuation at the back of the ready (round-robin scheduling) and returns the current kernel state. The scheduler uses this new kernel state to evaluate the next ready thread. The `yield` action using a continuation style is cumbersome, because one has to follow the continuation along and one has to perform an explicit `yield` often. The `yield` and thread-combinators resembling a more common monadic style. Our `return`, `>=>` and `>>` functions resemble the monadic functions of Haskell¹¹. Whenever a running thread performs an action such as a `return`, control is voluntarily given to the scheduler.

```

-- Thread a
a -> yield (cont x) k

-- Thread b
b = (a -> Thread b) -> Thread b
c -> l (\x -> r x cont) k

-- Thread r
r -> r

-- Main thread
main = newThread (print ['h', 'e', 'l', 'l', 'o']) >>
      print ['w', 'o', 'r', 'l', 'd']

-- End of main thread
return Void
printChar c >> print cs

```

The example above starts a thread that prints "hello" concurrently with a thread that prints "world". It assumes a low-level print function that prints a single character. The output of both threads is "hello", and is printed as "hweolrlloed".

Threads and Signals

(e.g. `newThread`) may fail because of external conditions such as other threads or operating system errors. Robust programs

elden and Rinus Plasmeijer

mented using dynamics. This makes it possible to store any
 n and to easily extend the set of exceptions at compile-time
 To provide this kind of exception handling, we extend the
 continuation argument for the case that an exception is

```
ExcCnt a) -> ExcCnt -> KernelOp
-> ExcCnt -> KernelOp
exception -> KernelOp
```

dynamic

```
and a | TC e
-> ec (dynamic e :: e~) k
```

```
on -> Thread a
= \sc ec k -> ec exception k
```

```
Exception -> Thread a) -> Thread a
=
and (\x _ -> sc x ec) (\e -> catcher e sc ec) k
```

ion wraps a value in a dynamic (hence the TC context re-
 s it to the enclosing **try** clause by evaluating the exception
throw can be used to throw an exception without wrapping
 in. The **try** function catches exceptions that occur during
 first argument (**thread**) and feeds it to its second argument
 any value can be thrown, exception handlers must match
 the exception using dynamic type pattern matching.
 des an outermost exception handler (not shown here) that
 hen an exception remains uncaught. This exception handler
 mer that an exception was not caught by any of the handlers
 of the occurring exception.

```
ead a
k -> yield (sc x ec) k
```

```
(a -> Thread b) -> Thread b
k -> l (\x -> r x sc) ec k
```



```
try (divide 42 0) handler
```

```
DivByZero
```

```
λ (x / y)
```

```
:: ArithErrors) = return 0 // or any other value
                = rethrow other
```

tion in the example throws the value `DivByZero` as an exception when the user tries to divide by zero. Exceptions caught in the body are handled by `handler`, which returns zero on a `DivByZero` exception. Exceptions of any other type are thrown again outside the `try`,

for concurrent setting, there is also a need for throwing and catching exceptions between different threads. We call this kind of inter-thread exceptions *signals*, or *asynchronous exceptions* as they are also called, as described by Marlow et. al. in an extension of Concurrent Haskell. We summarize our interface for signals below.

```
λ e -> Thread Void | TC e
```

```
λ a) -> Thread a
```

```
λ a) -> Thread a
```

Control is transferred from one thread to the other by the scheduler. A signal is thrown again when it arrives at the designated thread, and can be caught in the same way as other exceptions. To prevent interruption of an operation, one can enclose operations in a `signalsOff` clause, during which the thread is uninterruptible. Regardless of any nesting, `signalsOn` means interruptible and `signalsOff` always means non-interruptible. It is clear whether program code can or cannot be interrupted. The composition and nesting of program fragments that use these signals is clear: when a signal is caught, control goes to the exception handler and

elden and Rinus Plasmeijer

n Famke

will show how a programmer can execute groups of threads on multiple workstations, to construct distributed programs in

Microsoft Windows processes to provide preemptive task switching of threads running inside different processes. Once processes on one or more computers, threads can be started in any one to produce Famke's message passing primitives for communications and processes. The dynamic linker plays an essential role of a thread from one process to another.

Thread Communication

Safe communication between threads are Concurrent Haskell and Concurrent Clean's lazy graph copying [7].

ML-Vars do not scale very well to a distributed setting because described by Stolz and Huch in [12]. The first problem is that distributed garbage collection because they are first class objects, distributed or mobile setting. The second problem is that the error is generally unknown, which complicates reasoning about the failure of failing or moving processes. Automatic lazy graph copying work on objects that are distributed over multiple (remote) from the same two problems.

Cell [13,12] solves the problem by implementing an asynchronous message passing system using ports. Famke uses the same kind of channels are channels that vanish as soon as they are closed by the process containing the creating thread dies. Accessing a channel throws an exception. Using ports as the means of communication, there a port resides (at the process of the creating thread) and (explicitly or because the process died). In contrast with Clean we do not limit ports to a single reader (which could be achieved using Clean's uniqueness typing). The single reader restriction is that the port vanishes when the reader vanishes but we do not in practice.

abstract port id

```

just s -> case stringToDynamic s of
    (True, (m :: msg^)) -> return m
    other -> throw InvalidMessageAtPort
Nothing -> readPort port // make it appear blocking

portId msg) String -> Thread Void | TC msg
string -> Thread (PortId msg)      | TC msg

Dynamic -> String
String -> (Bool, Dynamic)

```

Ports operate on typed messages. The `newPort` function and `closePort` removes a port. `writePort` and `readPort` send and receive messages. The dynamic run-time system is used to send to and from a dynamic. Because we do not want to read a message until we want to send a message to someone, we will use the `String` and `stringToDynamic` functions from the dynamic library. These functions are similar to Haskell's `show` and `read`, `serialize` and `deserialize` functions and closures. They should be handled so they allow you to distinguish between objects that should be distinguished (e.g. between a closure and its value). The actual sending and receiving is done via simple message (string) passing primitives in the operating system. The `registerPort` function associates a port, by which the port can be looked up using `lookupPort`. In the untyped Haskell and Famke both use ports, our system is capable of sending and receiving functions (and therefore also closures) using the `registerPort` function. The dynamic type system also allows programs to refer to data of type `(PortId Dynamic)`, previously unknown data structures used by polymorphic functions or functions that work on the dynamic. The `dynamicApply` functions in section 2. An asynchronous system, such as presented here, allows programmers to build in synchronization methods (e.g. remote procedure calls, channels).

An example of a database server that uses a port to receive requests and applies them to the database.

```
list of records or something like that
```

elden and Rinus Plasmeijer

```
MyDBase" >>= \port ->
port mutateDatabase

DBase -> DBase
= ... // change the database
```

ces, and registers, a port that receives functions of type clients send functions that perform changes to the database rt. The server then waits for functions to arrive and applies e db. These functions can be safely applied to the database c run-time system guarantees that both the server and the e notion of the type of the database (**DBase**), even if they ograms. This is also an example of a running program that ded with new code.

Management

shows does the preemptive scheduling of processes, our sched- ny knowledge about multiple processes. Instead of changing t our system automatically add an additional thread, called *read*, to each process when it is created. This management ndle signals from other processes and to route them to the On request from threads running at other processes, it also a of new threads inside its own process. This management ion with the scheduler and the port implementation, form t is included in each process.

```
// abstract process id
ring

a -> Thread ProcId
Id (Thread a) -> Thread ThreadId
```

unction creates a new process at a given location and re- The creation of a new process is implemented by starting n executable, the *loader*, which becomes the new process. ple Clean program that starts a management thread. The on starts a new thread in another process. The thread is

```

newThreadAt pid (thread >>= writePort port) >>
return (Remote port)

-> Thread a | TC a
= readPort port >>= \result ->
  closePort port >>
  return result

```

tion creates a port to which the result of the given thread
 en starts a child thread at the remote location `pid` that
 and writes it to the port, and returns the port enclosed in
 e parent. When the parent decides that it wants the result,
 et it and to close the port.

Our system with this kind of heavyweight process enables
 build distributed concurrent applications. If one wants to run
 t contain parallel algorithms on a farm of workstations, this
 ver, non-trivial changes are required to the original program
 this. These changes include splitting the program code into
 d making communication between the threads explicit. The
 es is unfortunate, but our system was primarily designed for
 programs (and eventually mobile programs), not to speedup
 y running them on multiple processors.

Our discussion of the micro kernel and its interface that pro-
 reads (with exceptions and signals), processes and type-safe
 lues of any type between them. Now it is time to present the
 t makes use of these strongly typed concurrency primitives.

with Famke: The Shell

Introduce our shell that enables programmers to construct
 ograms interactively.

a way to interact with an operating system, usually via a
 ne/console interface. Normally, a shell does not provide a
 ing language, but it does enable users to start pre-compiled
 most shells provide simple ways to combine multiple pro-
 g and concurrent execution, and support execution-flow con-

elden and Rinus Plasmeijer

ample, it could test if a printing program (`:: WordDocument`
atches a document (`:: WordDocument`).
ts function application, variables, and a subset of Clean's
s. The shell syntax closely resembles Haskell's do-notation,
tions to read and write files.
e command line examples with an explanation of how they
shell.

0]

and `add` are unbound (do not appear in the left hand side
expression) in this example and our shell therefore assumes
of files (dynamics on disk). All files are supposed to contain
together represent a typed file system. The shell reads them
ically extending its functionality with these functions, and
f the dynamics. It uses the types of `map` (let us assume that
the type that we expect: $(a \rightarrow b) [a] \rightarrow [b]$), `add` (let us
e `Int`) and the list comprehension (which has type: `[Int]`)
command line. If this succeeds, which it should given the types
lies the partial application of `add` with the integer one to
rom one to ten, using the `map` function. The application of
ther is done using the `dynamicApply` function from Section
etter error reporting. With the help of the `dynamicApply`
constructs a new function that performs the computation `map`
This function uses the compiled code of `map`, `add`, and the list
shell is a hybrid interpreter/compiler, where the command
compiled to a function that is almost as efficient as the same
ectly in Clean and compiled to native code. Dynamics are
ting the command line, so it is not possible to change the
f the command line by overwriting a dynamic.

```
b inc [2,4..10]
```

le with the name `inc` as the partial application of the `add`
ger one. Then it applies the `map` function using the variable
n integers from two to ten. The dynamic linker detects that
eady linked in, and reuses their code.

```
b inc ['a'..'z']
```

ble `inc` as in the previous example, but applies it, using the

functions, which are read from disk by the shell before typing the command line, `result` is read in during the execution of `do`.

```
MyDBase"; writePort p (insertDBase MyRecord)
```

In the example above creates a new thread that executes the function `insertDBase`. Let us assume that we have two dynamics on disk: `insertDBase` containing a function that can insert a record and one with the name `MyRecord` containing a record for the second line, we get the port of the server by looking it up using `lookupPort`. We send the function `insertDBase` applied to `MyRecord` to the port. This example shows how we can communicate with threads in a type safe way.

Work

Both versions of both Haskell and Clean. Concurrent Haskell runs threads in a single UNIX process and provides M-Vars as communication between threads. Concurrent Clean [7] is only available on Transputers and on a network of single-processor Apple computers. Concurrent Clean provides support for native threads on transputers. On a network of Apple computers, it runs the same as on transputers, providing a virtual multiprocessor system. Concurrent Clean provided lazy graph copying as the primary communication mechanism. Concurrent systems cannot easily provide type safety between threads or between multiple incarnations of a single program.

The difference between Famke and the concurrent versions of Haskell and Clean is the lack of communication primitives. Neither lazy graph copying nor M-Vars can be used in a distributed setting because they require distributed memory. This issue has led to a distributed version of Concurrent Clean that does not use ports. However, its implementation does not allow functions to be sent over ports, because it cannot serialize functions. This problem could be provided by a dynamic linker for Concurrent Haskell. Lin and Lin [15] have extended Standard ML with threads (implemented as $M(\text{GC})$).

elden and Rinus Plasmeijer

built two prototypes of a Java operating system. Although its extensibility, portable byte code and static/dynamic type may be a way to build an operating system where multiple Java programs run concurrently, Java lacks the power of polymorphic and closures (to allow laziness) that our functional ap-

proach handles exception handling, while remaining pure and lazy. In [11] monadic exceptions has been added to Concurrent Haskell. Our approach closely follows their approach.

Es [18] integrates a shell into the programming language in order to allow the user to use the full expressiveness of Scheme. Es [19] is a higher-order language with first-class functions and allows the user to construct new commands on the command line. Neither shell provides a way to read and write files and to disk, and they cannot provide type safety because they are untyped executables.

Conclusions and Future Work

We presented the basics of our prototype functional operating system, Famke. Famke is written entirely in Clean and provides lightweight and heavyweight processes, and a type safe communication library. Clean's dynamic type system and dynamic linking support. We have built an interactive shell that type checks the command line and executes it. With the help of these mechanisms it becomes feasible to run concurrent Clean programs running on a network. Programs can be updated with new code at run-time using the dynamic run-time

approach to our kernel in a modular way by putting all extensions in modules which would allow us to tailor our system (at run-time) to the needs of the user. Nevertheless, there remain issues that need further research.

We would like to give the programmer more information about what exceptions are thrown. Unfortunately, we have not yet found a way to do this without compromising the flexibility of our approach.

The generation of ports given in this paper does not check if the name is already existing or even exists (when looking up), entrusting this to the programmer. Fortunately, this situation will be detected

and M. C. J. D. van Eekelen. *Functional Programming and Parallelism*. Addison Wesley, 1993.

and R. Plasmeijer. *Concurrent CLEAN Language Report (version 1.0)*. University of Nijmegen, December 2001. <http://www.cs.kun.nl/~clean/>

Adelli, B. Pierce, and G. Plotkin. Dynamic Typing in a Statically Typed Language. *ACM Transactions on Programming Languages and Systems*, 14(4):601–640, April 1991.

Arts, T. Types and Type Dependent Functions. In T. Davie K. Hammond and R. Plasmeijer, editors, *Proceedings of the 10th International Workshop on the Implementation of Functional Languages*, volume 1595 of *Lecture Notes in Computer Science*, pages 141–188. Springer-Verlag, 1998.

Arts, T. R. Plasmeijer. Lazy Dynamic Input/Output in the Lazy Functional Language Clean. In R. Peña and T. Arts, editors, *Proceedings of the 14th International Workshop on the Implementation of Functional Languages*, Lecture Notes in Computer Science. Springer-Verlag, 2002.

Arts, T., J.E.W. Smetsers, M.C.J.D. van Eekelen, and M.J. Plasmeijer. Lazy Dynamic Input/Output in the Lazy Functional Language Clean. In E.H.L. Aarts, J. van Leeuwen, and M. Rem, editors, *PARLE 1991: Architectures and Languages Europe, Volume II*, volume 506 of *Lecture Notes in Computer Science*, pages 202–219. Springer, 1991.

Arts, T. Evaluation-Based Multiprocessing. In J. Allen, editor, *Conference Record of the 1980 LISP Conference*, pages 19–28, Palo Alto, CA, 1980. The Lisp Society.

Barman, A. Poor Man’s Concurrency Monad. *Journal of Functional Programming*, 10(1):1–20, 2000.

Barman, A., A. Gordon, and S. Finne. Concurrent Haskell. In *Conference Record of the 1996: The 23rd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pages 295–308, St. Petersburg Beach, Florida, 21–24 January 1996. ACM Press.

Barman, A., J. Peyton Jones, A. Moran, and J.H. Reppy. Asynchronous Exception Handling in Haskell. In *SIGPLAN Conference on Programming Language Design and Implementation*, pages 274–285, 2001.

Barman, A. and J. Peyton Jones. Implementation of Port-based Distributed Haskell, 2001. <http://www.informatik.rwth-aachen.de/Research/distributedHaskell/>

Barman, A. and J. Peyton Jones. Distributed Programming in Haskell with Ports. In M. H. Kim and M. Koopman, editors, *Implementation of Functional Languages, Proceedings of the 1999 Workshop, IFL 2000*, volume 2011 of *Lecture Notes in Computer Science*, pages 107–121. Springer, September 2000.